

# Assignment

## Event Processing API

### Problem

**Build a backend service that implements an endpoint for transforming and storing event data, and another endpoint for reading the transformed output as JSON.**

#### Endpoint 1: Transform and Store Event Data

Implement an HTTP endpoint that allows for negotiation of some mechanism for the client to stream data to the server (e.g. HTTP long polling or a WebSocket). The endpoint should accept the **session ID** as a parameter. The format of the ID is up to you. Once the stream is open, the server should accept messages from the client in the following JSON format:

```
[
  {
    "timestamp": 1569972082,
    "type": "SESSION_START",
    "session_id": "e6fa79ca-d142-4046-a134-5134f16a0b5e",
  },
  {
    "timestamp": 1569972084,
    "type": "EVENT",
    "name": "cart_load"
  },
  ...
]
```

- Each message is a **JSON array** containing 1 or more JSON objects.
- Each **JSON object** contained within the array represents a single event.
- The client can send an unlimited number of messages containing 1 or more event entries.
- All objects contain a `type` key. The permissible entry types are `SESSION_START`, `SESSION_END` and `EVENT`. `SESSION_{START,END}` mark the start and end of the session. An `EVENT` is a single event occurrence that is recorded as a point with a timestamp.

- All entries contain a `timestamp` key, which is UNIX epoch time in seconds. Realistically, we would want higher precision (at least milliseconds), but we will use seconds here to simplify the problem.
- *Some* entries contain a `session_id` key. This is an identifier that is used to match related entries. For example, a `SESSION_START` and a `SESSION_END` entry pair will have the same session ID. In the above example, we used a `uuid4` format for the session ID, but this can be anything you choose.
- Some entry types contain additional keys that are specific to that event type. The documentation below on the different event types goes into more detail on this.
- **There is no ordering guarantee for entries**, except that `SESSION_START` will always be the first event, and `SESSION_END` will always be the last event. You could have events with timestamps that are not in chronological order.
- Once the `SESSION_END` event is received, the stream can be closed and the session is considered complete.

## Trace Events

### SESSION\_START

```
{
  "timestamp": 1569972082,
  "type": "SESSION_START",
  "session_id": "e6fa79ca-d142-4046-a134-5134f16a0b5e",
}
```

- There are no additional entry-specific keys.
- `session_id` for the matching `SESSION_END` entry will be the same.
- Guaranteed to be the **first** entry in the session.

### SESSION\_END

```
{
  "timestamp": 1569972090,
  "type": "SESSION_END",
  "session_id": "e6fa79ca-d142-4046-a134-5134f16a0b5e",
}
```

- There are no additional entry-specific keys.
- `session_id` for the matching `SESSION_START` entry will be the same.
- Guaranteed to be the **last** entry in the trace.

## EVENT

```
{
  "timestamp": 1569972083,
  "type": "EVENT",
  "name": "cart_loaded",
}
```

- There is one additional key, `name`, which is the name of the event.
- There is no `session_id`, since the event is just a single point.

## Data Transformation

The service should transform the raw trace data into a form that is more readily consumable by a front-end. **You do not have to implement a front-end.** This transformed data representation has the following key criteria:

- Events are ordered in **chronological** order.
- Each “pair” of related events (`SESSION_{START,END}`) is grouped and represented using a single object with `start` and `end` timestamps, and optionally `children` representing events that belong to a parent -- the parent is the session, which is the root object.

## Example

Let's take the following example of a complete input (which would have been delivered over the course of multiple streamed messages):

```
[
  {
    "timestamp": 1569972082,
    "type": "SESSION_START",
    "session_id": "e6fa79ca-d142-4046-a134-5134f16a0b5e",
  },
  {
    "timestamp": 1569972083,
    "type": "EVENT",
    "name": "cart_loaded"
  }
]
```

```

    },
    {
      "timestamp": 1569972090,
      "type": "SESSION_END",
      "session_id": "e6fa79ca-d142-4046-a134-5134f16a0b5e",
    }
  ]

```

- The entries start with `SESSION_START` and end with `SESSION_END`, as defined by the spec.
- There is a single `EVENT`.

The transformed output should look like this:

```

{
  "type": "SESSION",
  "start": 1569972082,
  "end": 1569972090,
  "children": [
    {
      "type": "EVENT",
      "timestamp": 1569972083,
      "name": "cart_loaded"
    }
  ]
}


```

The events are ordered, grouped, and represented in a hierarchical structure.

#### Edge Cases

- There's an edge case where you might have multiple events with identical timestamps. In this scenario the sorting doesn't matter.
- There may be other edge cases not discussed here. If you find any, make sure to note them and document how they would be handled.

#### Data Storage



The transformed output is persisted and should be retrievable by the **session ID** that was sent when the client first started the stream of events. There is no requirement on *how* or *where* the data needs to be stored -- that's up to you.

## Endpoint 2: Retrieve Transformed Data

Implement a second HTTP endpoint where you pass the **session ID** as a parameter, and the returned response is the transformed JSON representation of the session described above. Error cases should be handled appropriately (e.g. if there is no session matching the specified ID).

## Other Requirements

- Use git for version control, the git history should be included with your project submission (i.e. the `.git` directory should be intact). Avoid putting the entire project in a single commit, ideally your git history should show the evolution of the project.
- Include any instructions necessary to set up/bootstrap the service.

**There are *no constraints* on the resources you are allowed to use (feel free to search for anything you need) to solve this problem or how much time you can take.**

## Submission

Upon completion, please email your zipped project directory.